




Top row: CSI flatters a dam-break model applied to rainfall floods, while both physical checks catch it. Bottom row: the pattern reverses - CSI correctly reports the failure, but the rainfall-trained model is so wet on its own data (12.5x) that both the water balance and the timing look better on the flood type it never saw. Reporting any one of these three alone would give the wrong answer in one direction or the other. Single checkpoint (seed 0); not seed-replicated.